SQL: Transactions

Introduction to Databases

CompSci 316 Fall 2016

DUKE
COMPUTER SCIENCE

Transactions

- A transaction is a sequence of database operations with the following properties (ACID):
  - Atomic: Operations of a transaction are executed all-or-nothing, and are never left "half-done"
  - Consistency: Assume all database constraints are satisfied at the start of a transaction, they should remain satisfied at the end of the transaction
  - Isolation: Transactions must behave as if they were executed in complete isolation from each other
  - Durability: If the DBMS crashes after a transaction commits, all effects of the transaction must remain in the database when DBMS comes back up

SQL transactions

- A transaction is automatically started when a user executes an SQL statement
- Subsequent statements in the same session are executed as part of this transaction
  - Statements see changes made by earlier ones in the same transaction
  - Statements in other concurrently running transactions do not
- COMMIT command commits the transaction
  - Its effects are made final and visible to subsequent transactions
- ROLLBACK command aborts the transaction
  - Its effects are undone

Fine prints

- Schema operations (e.g., CREATE TABLE) implicitly commit the current transaction
  - Because it is often difficult to undo a schema operation
- Many DBMS support an AUTOCOMMIT feature, which automatically commits every single statement
  - You can turn it on/off through the API
    - Examples later in this lecture
  - For PostgreSQL:
    - `psql` command-line processor turns it on by default
    - You can turn it off at the `psql` prompt by typing:
      `\set AUTOCOMMIT 'off'`

Atomicity

- Partial effects of a transaction must be undone when
  - User explicitly aborts the transaction using ROLLBACK
    - E.g., application asks for user confirmation in the last step and issues COMMIT or ROLLBACK depending on the response
  - The DBMS crashes before a transaction commits

- Partial effects of a modification statement must be undone when any constraint is violated
  - Some systems roll back only this statement and let the transaction continue; others roll back the whole transaction

- How is atomicity achieved?
  - Logging (to support undo)

Durability

- DBMS accesses data on stable storage by bringing data into memory

- Effects of committed transactions must survive DBMS crashes

- How is durability achieved?
  - Forcing all changes to disk at the end of every transaction?
    - Too expensive
  - Logging (to support redo)

Consistency

- Consistency of the database is guaranteed by constraints and triggers declared in the database and/or transactions themselves
  - Whenever inconsistency arises, abort the statement or transaction, or (with deferred constraint checking or application-enforced constraints) fix the inconsistency within the transaction

Isolation

- Transactions must appear to be executed in a serial schedule (with no interleaving operations)

- For performance, DBMS executes transactions using a serializable schedule
  - In this schedule, operations from different transactions can interleave and execute concurrently
  - But the schedule is guaranteed to produce the same effects as a serial schedule

- How is isolation achieved?
  - Locking, multi-version concurrency control, etc.

SQL isolation levels

- Strongest isolation level: SERIALIZABLE
  - Complete isolation

- Weaker isolation levels: REPEATABLE READ, READ COMMITTED, READ UNCOMMITTED
  - Increase performance by eliminating overhead and allowing higher degrees of concurrency
  - Trade-off: sometimes you get the "wrong" answer

READ UNCOMMITTED

- Can read "dirty" data
  - A data item is dirty if it is written by an uncommitted transaction

- Problem: What if the transaction that wrote the dirty data eventually aborts?

- Example: wrong average
  -
```
-- T1:                      -- T2:
UPDATE User
SET pop = 0.99
WHERE uid = 142;            SELECT AVG(pop)
                            FROM User;

ROLLBACK;

                            COMMIT;
```

READ COMMITTED

- No dirty reads, but non-repeatable reads possible
  - Reading the same data item twice can produce different results
- Example: different averages
  -
```
-- T1:



UPDATE User
SET pop = 0.99
WHERE uid = 142;
COMMIT;
```
```
-- T2:
SELECT AVG(pop)
FROM User;




SELECT AVG(pop)
FROM User;
COMMIT;
```

REPEATABLE READ

- Reads are repeatable, but may see phantoms
- Example: different average (still!)

```
-- T1:                        -- T2:
                              SELECT AVG(pop)
                              FROM User;

INSERT INTO User
VALUES(789, 'Nelson',
       10, 0.1);
COMMIT;

                              SELECT AVG(pop)
                              FROM User;
                              COMMIT;
```

Summary of SQL isolation levels

Isolation level/anomaly	Dirty reads	Non-repeatable reads	Phantoms
READ UNCOMMITTED	Possible	Possible	Possible
READ COMMITTED	Impossible	Possible	Possible
REPEATABLE READ	Impossible	Impossible	Possible
SERIALIZABLE	Impossible	Impossible	Impossible

- Syntax: At the beginning of a transaction,
  SET TRANSACTION ISOLATION LEVEL
  isolation_level [READ ONLY|READ WRITE];
  - READ UNCOMMITTED can only be READ ONLY

- PostgreSQL defaults to READ COMMITTED

Transactions in programming

Using `pyscopg2` as an example:

```
conn = psycopg2.connect(dbname='beers')
conn.set_session(isolation_level='SERIALIZABLE',
                 ready_only=False,
                 autocommit=True)
```

- `isolation_level` defaults to READ COMMITTED
- `read_only` defaults to `False`
- `autocommit` defaults to `False`

- When `autocommit` is `False`, commit/abort current transaction as follows:

```
conn.commit()
```

```
conn.rollback()
```

ANSI isolation levels are lock-based

- `READ UNCOMMITTED`
  - Short-duration locks: lock, access, release immediately
- `READ COMMITTED`
  - Long-duration write locks: do not release write locks until commit
- `REPEATABLE READ`
  - Long-duration locks on all data items accessed
- `SERIALIZABLE`
  - Lock ranges to prevent insertion as well

Isolation levels not based on locks?

Snapshot isolation in Oracle

- Based on multiversion concurrency control
  - Used in Oracle, PostgreSQL, MS SQL Server, etc.
- How it works
  - Transaction X performs its operations on a private snapshot of the database taken at the start of X
  - X can commit only if it does not write any data that has been also written by a transaction committed after the start of X
- Avoids all ANSI anomalies
- But is NOT equivalent to SERIALIZABLE because of write skew anomaly

Write skew example

- Constraint: combined balance $A + B \geq 0$

- $A = 100, B = 100$

- T_1 checks $A + B - 200 \geq 0$, and then proceeds to withdraw 200 from A

- T_2 checks $A + B - 200 \geq 0$, and then proceeds to withdraw 200 from B

- Possible under snapshot isolation because the writes (to A and to B) do not conflict

- But $A + B = -200 < 0$ afterwards!

Bottom line

- Group reads and dependent writes into a transaction in your applications
  - E.g., enrolling a class, booking a ticket

- Anything less than `SERIALABLE` is potentially very dangerous
  - Use only when performance is critical
  - `READ ONLY` makes weaker isolation levels a bit safer